

## USDA LLM Model Pipeline

### 1. What the pipeline does (high level)

Stage	Script	What it Produces
Harvest & clean	API_Document_CommentsDownloader.py	One or more JSON files, each containing 'comment_id', full text, and a list of local PDF/Image attachments. All attachments are saved to <i>./Downloads/USDA_JSON/attachments</i> by default.
Summarize with GPT4o	LLM_Model_FINAL.py	Processed_comments.csv with columns: who_type, who_name, what, why, issues, scientific legal support, etc
Bucket issues	Comment_clustering.py	processed_with_categories.csv (each comment now has <b>high-level issue categories</b> ) sorted_by_issue.csv (one row per comment × category pair, ready for pivoting) Two log files capturing every GPT response (for audit)

NOTE: The output of each stage is the required input of the next, so you **run the scripts in the above order**.

### 2. Environment & One-Time Setup

All three scripts are pure Python 3.10+. You can keep them in the same project folder.

#### Install system tools

- Windows:
  - [Tesseract-OCR](#) → default path C:\Program Files\Tesseract-OCR (used for OCR).
  - Poppler – place poppler/bin on your PATH (used by pdf2image).

Install below packages using -

```
pip install requests pandas pdfplumber pdf2image pytesseract pillow \
python-dotenv openai
```

Create a .env in the project root –

OPENAI\_API\_KEY=sk-xx

(The loader is called in the LLM and clustering scripts)

### 3. Stage-by-stage walkthrough

#### 3.1 API\_Document\_CommentsDownloader.py

1. **User supplies a Regulations.gov URL** (document or docket).  
The helper `extract_docket_id_or_document_id()` parses the URL and classifies it .
2. **Find the internal objectId** for that document via the Regulations.gov REST API (`/v4/documents/{id}`) .
3. **Iterate through comments (250 per page)**
  - For each comment the script hits the comments endpoint, then follows the comment's **self link** to get the full body text .
  - Attachments are listed via `/comments/{comment_id}/attachments`.  
Every file is streamed to disk in `./Downloads/USDA_JSON/attachments`, preserving its original MIME type.
4. **Light cleaning & incremental saving**  
Only three keys are stored: `comment_id`, `text`, `attachments`.  
Each page is written to its own file (e.g. `comments_FSYS-2023-0001_page_3.json`) so partial runs can be resumed easily.  
→ *Output ready for Stage 2.*

**Change API key or rate:** edit the `API_KEY` constant and/or `sleep(0.1)` delay between comment downloads.

#### 3.2 LLM\_Model\_FINAL.py – comment summarization

1. **Configuration block (top of the file)**
  - `JSON_FOLDER`, `PDF_FOLDER`, `OUTPUT_FILE` – point to Stage 1 outputs.
  - `USE_API=False` lets you dry-run without burning tokens .
  - `MAX_TOKENS` is a hard truncate safeguard before sending text to GPT-4-turbo.
2. **Reading comments & attachments**

- For every JSON page file, each comment's body text is taken.
  - PDFs are opened locally; text is extracted by **pdfplumber** first, and (if empty) by **Tesseract OCR** on page images .
3. **Prompt-driven extraction**
- The `classify_comment_by_issue()` function contains the full system & user prompt ( $\approx 80$  lines). It asks GPT-4-turbo to return valid JSON with keys:  
`who_type`, `who_name`, `what`, `why`, `issues[]`, `scientific_legal_support` .
4. **Error handling**
- Regex `extract_json_block()` strips Markdown fences and pulls out the first JSON object .
  - If parsing fails, a safe fallback row is inserted and the comment is flagged for manual review.
5. **Export**
- A single processed `comments.csv` is written with one row per comment.  
 → *Output ready for Stage 3.*

### How to edit the prompt

- \*Modify the big triple-quoted prompt variable inside `classify_comment_by_issue()`.
- \*Change the model (`model="gpt-4-turbo"`) or temperature (defaults to 0).

## 3.3 `comment_clustering.py` – thematic bucketing

1. **Load the CSV from Stage 2** (`INPUT_FILE` constant).  
 The issues column is split into Python lists.
2. **Batch set of unique issues** (`Counter > all_issues`).  
 Large dockets may have  $>1\ 000$  unique phrases, so issues are chunked into  $\leq 500$  per GPT call .
3. **First GPT task – “Group issues into 8–15 categories”**
  - Prompt builder `build_prompt()` shows *every* issue as a bullet list; GPT must return an array of `{category, related_issues[]}` objects.
  - Each raw response is saved to `gpt_issue_grouping_raw.txt` for transparency.
4. **Second GPT task – “Merge near-duplicate category names”**  
 This guard-rail collapses e.g. “Worker Safety and Health” & “Worker Conditions” into a single canonical label using another JSON-returning prompt .
5. **Map every comment to one or more high-level categories** (`high_level_issues` list).

- `processed_with_categories.csv` keeps the original wide format.
- `sorted_by_issue.csv` explodes the list so each comment-category pair is a row (handy for Power BI filters or pivot tables).

### To change how clustering works

*Tweak* `build_prompt()` (lines near the top of the file) or the consolidation prompt in `consolidate_categories()`.

*Swap* the model to `gpt-4o`, `gpt-4o-mini`, etc.

*Adjust* `BATCH_SIZE` to trade off speed vs. token load.

### CUSTOMIZATION CHEAT SHEET:

Change you want	Where to edit
Use a different Regulations.gov API key	Top of <code>API_Document_CommentsDownloader.py</code> ( <code>API_KEY</code> )
Switch from per-document to per-docket harvesting	Accept docket IDs in <code>extract_docket_id_or_document_id()</code> and pass them to <code>get_comments_by_object_id()</code> (API supports both)
Store downloads somewhere else	Paths under <code>download_file()</code> and <code>JSON_FOLDER/PDF_FOLDER</code> constants
Try GPT-4o or reduce temperature	The model and temperature parameters in each <code>client.chat.completions.create()</code> call
Add / remove output columns	In Stage 2, edit the data dict before it is appended to results
Change the number of high-level categories	Edit the line “Return only 8–15 categories TOTAL” in <code>build_prompt()</code>